

Short papers

Authors	Title
Heike Adel and Hinrich Schütze	<i>Global Normalization of Convolutional Neural Networks for Joint Entity and Relation Classification</i>
Alham Fikri Aji and Kenneth Heafield	<i>Sparse Communication for Distributed Gradient Descent</i>
Md Shad Akhtar, Abhishek Kumar, Deepanway Ghosal, Asif Ekbal and Pushpak Bhattacharyya	<i>A Multilayer Perceptron based Ensemble Technique for Fine-grained Financial Sentiment Analysis</i>
Jacob Andreas and Dan Klein	<i>Analogs of Linguistic Structure in Deep Representations</i>
Tatsuya Aoki, Ryohei Sasano, Hiroya Takamura and Manadu Okumura	<i>Distinguishing Japanese Non-standard Usages from Standard Ones</i>
Zhenisbek Assylbekov, Rustem Takhanov, Bagdat Myrzakhmetov and Jonathan N. Washington	<i>Syllable-aware Neural Language Models: A Failure to Beat Character-aware Ones</i>
Miguel Ballesteros and Yaser Al-Onaizan	<i>AMR Parsing using Stack-LSTMs</i>
Antonio Valerio Miceli Barone, Barry Haddow, Ulrich Germann and Rico Sennrich	<i>Regularization techniques for fine-tuning in neural machine translation</i>
Rachel Bawden	<i>Machine Translation, it's a question of style, innit? The case of English tag questions</i>
Jan A. Botha, Emily Pitler, Ji Ma, Anton Bakalov, Alex Salcianu, David Weiss, Ryan McDonald and Slav Petrov	<i>Natural Language Processing with Small Feed-Forward Networks</i>
Jiong Cai, Yong Jiang and Kewei Tu	<i>CRF Autoencoder for Unsupervised Dependency Parsing</i>
Andrew Cattle and Xiaojuan Ma	<i>Predicting Word Association Strengths</i>

Authors	Title
Yin-Wen Chang and Michael Collins	<i>Source-Side Left-to-Right or Target-Side Left-to-Right? An Empirical Comparison of Two Phrase-Based Decoding Algorithms</i>
Kehai Chen, Rui Wang, Masao Utiyama, Lemao Liu, Akihiro Tamura, Eiichiro Sumita and Tiejun Zhao	<i>Neural Machine Translation with Source Dependency Representation</i>
Prafulla Kumar Choubey and Ruihong Huang	<i>A Sequential Model for Classifying Temporal Relations between Intra-Sentence Events</i>
Oana Cocarascu and Francesca Toni	<i>Identifying attack and support argumentative relations using deep learning</i>
Ryan Cotterell, Ekaterina Vylomova, Huda Khayrallah, Christo Kirov and David Yarowsky	<i>Paradigm Completion for Derivational Morphology</i>
Jacob Devlin	<i>Sharp Models on Dull Hardware: Fast and Accurate Neural Machine Translation Decoding on the CPU</i>
Tobias Domhan and Felix Hieber	<i>Using Target-side Monolingual Data for Neural Machine Translation through Multi-task Learning</i>
Zi-Yi Dou	<i>Capturing User and Product Information for Document Level Sentiment Analysis with Deep Memory Network</i>
Xinya Du and Claire Cardie	<i>Identifying Where to Focus in Reading Comprehension for Neural Question Generation</i>
Mihai Dusmanu, Elena Cabrio and Serena Villata	<i>Argument Mining on Twitter: Arguments, Facts and Sources</i>
Jared Fernandez, Zhaocheng Yu and Doug Downey	<i>VecShare: A Framework for Sharing Word Representation Vectors</i>
Annemarie Friedrich and Damyana Gateva	<i>Classification of telicity using cross-linguistic annotation projection</i>
Cheng-Yang Fu, Joon Lee, Mohit Bansal and Alexander Berg	<i>Video Highlight Prediction Using Audience Chat Reactions</i>

Authors	Title
Varun Gangal, Harsh Jhamtani, Graham Neubig, Eduard Hovy and Eric Nyberg	<i>Charmanteau: Character Embedding Models For Portmanteau Creation</i>
Clément Gautrais, Peggy Cellier, René Quiniou and Alexandre Termier	<i>Topic Signatures in Political Campaign Speeches</i>
Spandana Gella, Rico Sennrich, Frank Keller and Mirella Lapata	<i>Image Pivoting for Learning Multilingual Multimodal Representations</i>
Tanya Goyal, Sachin Kelkar, Manas Agarwal and Jeenu Grover	<i>An Empirical Analysis of Edit Importance between Document Versions</i>
Arihant Gupta, Syed Sarfaraz Akhtar, Avijit Vajpayee, Arjit Srivastava, Madan Gopal Jhanwar and Manish Shrivastava	<i>Exploiting Morphological Regularities in Distributional Word Representations</i>
E. Dario Gutierrez, Guillermo Cecchi, Cheryl Corcoran and Philip Corlett	<i>Using Automated Metaphor Identification to Aid in Detection and Prediction of First-Episode Schizophrenia</i>
Dan Han, Pascual Martínez-Gómez and Koji Mineshima	<i>Visual Denotations for Recognizing Textual Entailment</i>
Wenjuan Han, Yong Jiang and Kewei Tu	<i>Dependency Grammar Induction with Neural Lexicalization and Big Training Data</i>
Souleiman Hasan and Edward Curry	<i>Word Re-Embedding via Manifold Dimensionality Retention</i>
Benjamin Heinzerling, Nafise Sadat Moosavi and Michael Strube	<i>Revisiting Selectional Preferences for Coreference Resolution</i>
Aurélie Herbelot and Marco Baroni	<i>High-risk learning: acquiring new word vectors from tiny data</i>
YiYao Huang and William Yang Wang	<i>Deep Residual Learning for Weakly-Supervised Relation Extraction</i>
Liang Huang, Kai Zhao and Mingbo Ma	<i>When to Finish? Optimal Beam Search for Neural Text Generation (modulo beam size)</i>

Authors	Title
Zhuoxuan Jiang, Shanshan Feng, Gao Cong, Chunyan Miao and Xiaoming Li	<i>A Novel Cascade Model for Learning Latent Similarity from Heterogeneous Sequential Data of MOOC</i>
Yong Jiang, Wenjuan Han and Kewei Tu	<i>Combining Generative and Discriminative Approaches to Unsupervised Dependency Parsing via Dual Decomposition</i>
Khalid Al Khatib, Henning Wachsmuth, Matthias Hagen and Benno Stein	<i>Patterns of Argumentation Strategies across Topics</i>
Joo-Kyung Kim, Young-Bum Kim, Ruhi Sarikaya and Eric Fosler-Lussier	<i>Cross-Lingual Transfer Learning for POS Tagging without Cross-Lingual Resources</i>
Satwik Kottur, José Moura, Stefan Lee and Dhruv Batra	<i>Natural Language Does Not Emerge ‘Naturally’ in Multi-Agent Dialog</i>
Andrey Kutuzov, Erik Velldal and Lilja Øvrelid	<i>Temporal dynamics of semantic relations in word embeddings: an application to predicting armed conflict participants</i>
Thomas Lavergne and François Yvon	<i>Learning the Structure of Variable-Order CRFs: a finite-state perspective</i>
Xin Li and Wai Lam	<i>Deep Multi-Task Learning for Aspect Term Extraction with Memory Interaction</i>
Peng-Hsuan Li, Ruo-Ping Dong, Yu-Siang Wang, Ju-Chieh Chou and Wei-Yun Ma	<i>Leveraging Linguistic Structures for Named Entity Recognition with Bidirectional Recursive Neural Networks</i>
Shen Li, Zhe Zhao, Tao Liu, Renfen Hu and Xiaoyong Du	<i>Initializing Convolutional Filters with Semantic Features for Text Classification</i>
Haijing Liu, Yang Gao, Pin Lv, Mengxue Li, Shiqiang Geng, Minglan Li and Hao Wang	<i>Using Argument-based Features to Predict and Analyse Review Helpfulness</i>
Tianyu Liu, Kexiang Wang, Baobao Chang and Zhifang Sui	<i>A Soft-label Method for Noise-tolerant Distantly Supervised Relation Extraction</i>

Authors	Title
Sharid Loáiciga, Liane Guillou and Christian Hardmeier	<i>What is it? Disambiguating the different readings of the pronoun ‘it’</i>
You Lu, Jeffrey Lund and Jordan Boyd-Graber	<i>Why ADAGRAD Fails for Online Topic Modeling</i>
Chaitanya Malaviya, Graham Neubig and Patrick Littell	<i>Learning Language Representations for Typology Prediction</i>
Nitika Mathur, Timothy Baldwin and Trevor Cohn	<i>Sequence Effects in Crowdsourced Annotations</i>
Oren Melamud, Ido Dagan and Jacob Goldberger	<i>A Simple Language Model based on PMI Matrix Approximations</i>
Stefano Menini, Federico Nanni, Simone Paolo Ponzetto and Sara Tonelli	<i>Topic-Based Agreement and Disagreement in US Electoral Manifestos</i>
David Mimno and Laure Thompson	<i>The strange geometry of skip-gram with negative sampling</i>
Mathieu Morey, Philippe Muller and Nicholas Asher	<i>How much progress have we made on RST discourse parsing? A replication study of recent results on the RST-DT</i>
Jesse Mu, Joshua K. Hartshorne and Timothy O'Donnell	<i>Evaluating Hierarchies of Verb Argument Structure with Hierarchical Clustering</i>
Vivi Nastase and Carlo Strapparava	<i>Word Etymology as Native Language Interference</i>
Xing Niu, Marianna Martindale and Marine Carpuat	<i>A Study of Style in Machine Translation: Controlling the Formality of Machine Translation Output</i>
Enrique Noriega-Atala, Marco A. Valenzuela-Escárcega, Clayton Morrison and Mihai Surdeanu	<i>Learning what to read: Focused machine reading</i>
Takamasa Oshikiri	<i>Segmentation-Free Word Embedding for Unsegmented Languages</i>
Natalie Parde and Rodney Nielsen	<i>Finding Patterns in Noisy Crowds: Regression-based Annotation Aggregation for Crowdsourced Data</i>

Authors	Title
Ramakanth Pasunuru and Mohit Bansal	<i>Reinforcement Learning for Video Captioning with Entailment Rewards</i>
Nima Pourdamghani and Kevin Knight	<i>Deciphering Related Languages</i>
Daniel Preoȕiuc-Pietro, Sharath Chandra Guntuku and Lyle Ungar	<i>Controlling Human Perception of Basic User Traits</i>
Jay Pujara, Eriq Augustine and Lise Getoor	<i>Sparsity and Noise: Where Knowledge Graph Embeddings Fall Short</i>
Hannah Rashkin, Eunsol Choi, Jin Yea Jang, Svitlana Volkova and Yejin Choi	<i>Truth of Varying Shades: Analyzing Language in Fake News and Political Fact-Checking</i>
Joseph Sanu, Mingbin Xu, Hui Jiang and Quan Liu	<i>Word Embeddings based on Fixed-Size Ordinally Forgetting Encoding</i>
Maarten Sap, Marcella Cindy Prasetio, Ari Holtzman, Hannah Rashkin and Yejin Choi	<i>Connotation Frames of Power and Agency in Modern Films</i>
Kanthashree Mysore Sathyendra, Shomir Wilson, Florian Schaub, Sebastian Zimmeck and Norman Sadeh	<i>Identifying the Provision of Choices in Privacy Policy Text</i>
Allen Schmaltz, Yoon Kim, Alexander Rush and Stuart Shieber	<i>Adapting Sequence Models for Sentence Correction</i>
Raksha Sharma, Arpan Somani, Lakshya Kumar and Pushpak Bhattacharyya	<i>Sentiment Intensity Ranking among Adjectives Using Sentiment Bearing Word Embeddings</i>
Aditya Sharma, Zarana Parekh and Partha Talukdar	<i>Speeding up Reinforcement Learning-based Information Extraction Training using Asynchronous Methods</i>
Eyal Shnarch, Ran Levy, Vikas Raykar and Noam Slonim	<i>GRASP: Rich Patterns for Argumentation Mining</i>
Lei Shu, Hu Xu and Bing Liu	<i>DOC: Deep Open Classification of Text Documents</i>

Authors	Title
Daniil Sorokin and Iryna Gurevych	<i>Context-Aware Representations for Knowledge Base Relation Extraction</i>
Lucas Sterckx, Jason Naradowsky, Bill Byrne, Thomas Demeester and Chris Develder	<i>Break it Down for Me: A Study in Automated Lyric Annotation</i>
Mitchell Stern, Daniel Fried and Dan Klein	<i>Effective Inference for Generative Neural Parsing</i>
Karl Stratos	<i>A Sub-Character Architecture for Korean Language Processing</i>
Siliang Tang, Ning Zhang, Jinjiang Zhang, Fei Wu and Yueting Zhuang	<i>NITE: A Neural Inductive Teaching Framework for Domain Specific NER</i>
Quan Hung Tran, Ingrid Zukerman and Gholamreza Haffari	<i>Preserving Distributional Information in Dialogue Act Classification</i>
Ferhan Ture and Oliver Jojic	<i>No Need to Pay Attention: Simple Recurrent Neural Networks Work!</i>
Kateryna Tymoshenko, Daniele Bonadiman and Alessandro Moschitti	<i>Ranking Kernels for Structures and Embeddings: A Hybrid Preference and Classification Model</i>
Ashwin Vijayakumar, Ramakrishna Vedantam and Devi Parikh	<i>Sound-Word2Vec: Learning Word Representations Grounded in Sounds</i>
Longyue Wang, Zhaopeng Tu, Andy Way and Qun Liu	<i>Exploiting Cross-Sentence Context for Neural Machine Translation</i>
Rui Wang, Masao Utiyama, Lemao Liu, Kehai Chen and Eiichiro Sumita	<i>Instance Weighting for Neural Machine Translation Domain Adaptation</i>
Shaonan Wang, Jiajun Zhang and Chengqing Zong	<i>Exploiting Word Internal Structures for Generic Chinese Sentence Representation</i>
Liang Wang, Sujian Li, Yajuan Lv and Houfeng WANG	<i>Learning to Rank Semantic Coherence for Topic Segmentation</i>
Yi Wu, David Bamman and Stuart Russell	<i>Adversarial Training for Relation Extraction</i>

Authors	Title
Hainan Xu and Philipp Koehn	<i>Zipporah: a Fast and Scalable Data Cleaning System for Noisy Web-Crawled Parallel Corpora</i>
Min Yang, Jincheng Mei, Heng Ji, zhao wei, Zhou Zhao and Xiaojun Chen	<i>Identifying and Tracking Sentiments and Topics from Social Media Texts during Natural Disasters</i>
Wei Yang, Wei Lu and Vincent Zheng	<i>A Simple Regularization-based Algorithm for Learning Cross-Domain Word Embeddings</i>
Weiwei Yang, Jordan Boyd-Graber and Philip Resnik	<i>Adapting Topic Models using Lexical Associations with Tree Priors</i>
Semih Yavuz, Izzeddin Gur, Yu Su and Xifeng Yan	<i>Recovering Question Answering Errors via Query Revision</i>
Licheng Yu, Mohit Bansal and Tamara Berg	<i>Hierarchically-Attentive RNN for Album Summarization and Storytelling</i>
Liang-Chih Yu, Jin Wang, K. Robert Lai and Xuejie Zhang	<i>Refining Word Embeddings for Sentiment Analysis</i>
Jinxing Yu, Xun Jian, Hao Xin and Yangqiu Song	<i>Joint Embeddings of Chinese Words, Characters, and Fine-grained Subcharacter Components</i>
Sina Zarrieß and David Schlangen	<i>Deriving continous grounded meaning representations from referentially structured multimodal contexts</i>
Xiaowei Zhang, Wei Chen, Shuang Xu and Bo Xu	<i>Towards Compact and Fast Neural Machine Translation Using a Combined Method</i>
Qing Zhang and Houfeng Wang	<i>Noise-Clustered Distant Supervision for Relation Extraction: A Nonparametric Bayesian Perspective</i>
Zhirui Zhang, Shujie Liu, Mu Li, Ming Zhou and Enhong Chen	<i>Stack-based Multi-layer Attention for Transition-based Dependency Parsing</i>
Jianmin Zhang and Xiaojun Wan	<i>Towards Automatic Construction of News Overview Articles by News Synthesis</i>



Authors	Title
Kai Zhao, Dapeng Li and Liang Huang	<i>Joint Syntactic and Discourse Parsing with Recurrent Neural Models</i>
Jie Zhao, Yu Su, Ziyu Guan and Huan Sun	<i>An End-to-End Deep Framework for Answer Triggering with a Novel Group-Level Objective</i>
Hao Zhou, Zhenting Yu, Yue Zhang, Shujian Huang, XIN-YU DAI and Jiajun Chen	<i>Word-Context Character Embeddings for Chinese Word Segmentation</i>

Long papers

Authors	Title
David Alvarez-Melis and Tommi Jaakkola	<i>A causal framework for explaining the predictions of black-box sequence-to-sequence models</i>
Hadi Amiri, Timothy Miller and Guergana Savova	<i>Repeat before Forgetting: Spaced Repetition for Efficient and Effective Training of Neural Networks</i>
Peter Anderson, Basura Fernando, Mark Johnson and Stephen Gould	<i>Guided Open Vocabulary Image Captioning with Constrained Beam Search</i>
Yuki Arase and Jun'ichi Tsujii	<i>Monolingual Phrase Alignment on Parse Forests</i>
Adam St Arnaud, David Beck and Grzegorz Kondrak	<i>Identifying Cognate Sets Across Dictionaries of Related Languages</i>
Ehsaneddin Asgari and Hinrich Schütze	<i>Past, Present, Future: A Computational Investigation of the Typology of Tense in 1000 Languages</i>
Ivano Basile and Fabio Tamburini	<i>Towards Quantum Language Models</i>

Authors	Title
Joost Bastings, Ivan Titov, Wilker Aziz, Diego Marcheggiani and Khalil Simaan	<i>Graph Convolutional Encoders for Syntax-aware Neural Machine Translation</i>
Chloé Braud, Ophélie Lacroix and Anders Søgaard	<i>Does syntax help discourse segmentation? Not so much</i>
Denny Britz, Melody Guan and Minh-Thang Luong	<i>Efficient Attention using a Fixed-Size Memory Representation</i>
Denny Britz, Anna Goldie, Minh-Thang Luong and Quoc Le	<i>Exploring Hyperparameter Sensitivity in Neural Machine Translation Architectures</i>
Luana Bulat, Stephen Clark and Ekaterina Shutova	<i>Speaking, Seeing, Understanding: Correlating semantic models with conceptual representation in the brain</i>
Iacer Calixto, Qun Liu and Nick Campbell	<i>Incorporating Global Visual Features into Attention-based Neural Machine Translation.</i>
Junjie Cao, Sheng Huang, Weiwei Sun and Xiaojun Wan	<i>Quasi-Second-Order Parsing for 1-Endpoint-Crossing, Pagenumber-2 Graphs</i>
Arun Chaganty, Ashwin Paranjape, Percy Liang and Christopher D. Manning	<i>Importance sampling for unbiased on-demand evaluation of knowledge base population</i>
Cheng Chang, Runzhe Yang, Lu Chen, Xiang Zhou and Kai Yu	<i>Affordable On-line Dialogue Policy Learning</i>
Snigdha Chaturvedi, Haoruo Peng and Dan Roth	<i>Story Comprehension for Predicting What Happens Next</i>
Peng Chen, Zhongqian Sun, Lidong Bing and Wei Yang	<i>Recurrent Attention Network on Memory for Aspect Sentiment Analysis</i>
Lu Chen, Xiang Zhou, Cheng Chang, Runzhe Yang and Kai Yu	<i>Agent-Aware Dropout DQN for Safe and Efficient On-line Dialogue Policy Learning</i>
Hao Cheng, Hao Fang and Mari Ostendorf	<i>A Factored Neural Network Model for Characterizing Online Discussions in Vector Space</i>

Authors	Title
Prafulla Kumar Choubey and Ruihong Huang	<i>Event Coreference Resolution by Iteratively Unfolding Inter-dependencies among Events</i>
Alexis Conneau, Douwe Kiela, Holger Schwenk, Loïc Barrault and Antoine Bordes	<i>Supervised Learning of Universal Sentence Representations from Natural Language Inference Data</i>
Caio Corro, Joseph Le Roux and Mathieu Lacroix	<i>Efficient Discontinuous Phrase-Structure Parsing via the Generalized Maximum Spanning Arborescence</i>
Ryan Cotterell and Georg Heigold	<i>Cross-lingual Character-Level Neural Morphological Tagging</i>
Leonard Dahlmann, Evgeny Matusov, Pavel Petrushkov and Shahram Khadivi	<i>Neural Machine Translation Leveraging Phrase-based Models in a Hybrid Search</i>
Johannes Daxenberger, Steffen Eger, Ivan Habernal, Christian Stab and Iryna Gurevych	<i>What is the Essence of a Claim? Cross-Domain Claim Identification</i>
Jean-Benoit Delbrouck and Stéphane Dupont	<i>An empirical study of the effectiveness of images on Multi-modal Neural Machine Translation</i>
Tao Ding, Warren K. Bickel and Shimei Pan	<i>Multi-View Unsupervised User Feature Embedding for Social Media-based Substance Use Prediction</i>
Li Dong, Jonathan Mallinson, Siva Reddy and Mirella Lapata	<i>Learning to Paraphrase for Question Answering</i>
Nan Duan, Duyu Tang, Peng Chen and Ming Zhou	<i>Question Generation for Question Answering</i>
haim dubossarsky, Daphna Weinshall and Eitan Grossman	<i>Outta Control: Laws of Semantic Change and Inherent Biases in Word Representation Models</i>
Greg Durrett, Jonathan K. Kummerfeld, Taylor Berg-Kirkpatrick, Rebecca Portnoff, Sadia Afroz, Damon McCoy, Kirill Levchenko and Vern Paxson	<i>Identifying Products in Online Cybercrime Marketplaces: A Dataset for Fine-grained Domain Adaptation</i>

Authors	Title
Joshua Eisenberg and Mark Finlayson	<i>A Simpler and More Generalizable Story Detector using Verb and Character Features</i>
Micha Elsner and Cory Shain	<i>Speech segmentation with a neural encoder model of working memory</i>
Arash Eshghi, Igor Shalyminov and Oliver Lemon	<i>Bootstrapping incremental dialogue systems from minimal data: the generalisation power of dialogue grammars</i>
Tobias Falke and Iryna Gurevych	<i>Bringing Structure into Summaries: Crowdsourcing a Benchmark Corpus of Concept Maps</i>
Meng Fang, Yuan Li and Trevor Cohn	<i>Learning how to Active Learn: A Deep Reinforcement Learning Approach</i>
Bjarke Felbo, Alan Mislove, Anders Søgaard, Iyad Rahwan and Sune Lehmann	<i>Using millions of emoji occurrences to learn any-domain representations for detecting sentiment, emotion and sarcasm</i>
Yang Feng, Shiyue Zhang, Andi Zhang, Dong Wang and Andrew Abel	<i>Memory-augmented Neural Machine Translation</i>
Besnik Fetahu, Katja Markert and Avishek Anand	<i>Fine Grained Citation Span for References in Wikipedia</i>
Kathleen C. Fraser, Kristina Lundholm Fors, Dimitrios Kokkinakis and Arto Nordlund	<i>An analysis of eye-movements during reading for the detection of mild cognitive impairment</i>
Lea Frermann and György Szarvas	<i>Inducing Semantic Micro-Clusters from Deep Multi-View Representations of Novels</i>
Zhe Gan, Yunchen Pu, Ricardo Henao, Chunyuan Li, Xiaodong He and Lawrence Carin	<i>Learning Generic Sentence Representations Using Convolutional Neural Networks</i>
Octavian-Eugen Ganea and Thomas Hofmann	<i>Deep Joint Entity Disambiguation with Local Neural Attention</i>
Aparna Garimella, Carmen Banea and Rada Mihalcea	<i>Demographic-aware word associations</i>

Authors	Title
Kiril Gashteovski, Rainer Gemulla and Luciano Del Corro	<i>MinIE: Minimizing Facts in Open Information Extraction</i>
Aniruddha Ghosh and Tony Veale	<i>Magnets for Sarcasm: Making Sarcasm Detection Timely, Contextual and Very Personal</i>
Goran Glavaš and Simone Paolo Ponzetto	<i>Dual Tensor Model for Detecting Asymmetric Lexico-Semantic Relations</i>
David Golub, Po-Sen Huang, Xiaodong He and Li Deng	<i>Two-Stage Synthesis Networks for Transfer Learning in Machine Comprehension</i>
Chen Gong, Zhenghua Li, Min Zhang and Xinzhou Jiang	<i>Multi-Grained Chinese Word Segmentation</i>
Jiatao Gu, Kyunghyun Cho and Victor O.K. Li	<i>Trainable Greedy Decoding for Neural Machine Translation</i>
Lin Gui, Jiannan Hu, Yulan He, Ruifeng Xu, Lu Qin and Jiachen Du	<i>A Question Answering Approach for Emotion Cause Extraction</i>
Tao Gui, Qi Zhang, Haoran Huang, Minlong Peng and Xuanjing Huang	<i>Part-of-Speech Tagging for Twitter with Adversarial Neural Networks</i>
Nitish Gupta, Sameer Singh and Dan Roth	<i>Entity Linking via Joint Encoding of Types, Descriptions, and Context</i>
Kata Gábor, Haifa Zargayouna, Isabelle Tellier, Davide Buscaldi and Thierry Charnois	<i>Exploring Vector Spaces for Semantic Relations</i>
Kazuma Hashimoto and Yoshimasa Tsuruoka	<i>Neural Machine Translation with Source-Side Latent Graph Parsing</i>
Kazuma Hashimoto, caiming xiong, Yoshimasa Tsuruoka and Richard Socher	<i>A Joint Many-Task Model: Growing a Neural Network for Multiple NLP Tasks</i>
Hua He, Kris Ganjam, Navendu Jain, Jessica Lundin, Ryen White and Jimmy Lin	<i>An Insight Extraction System on BioMedical Literature with Deep Neural Networks</i>

Authors	Title
Daniel Hewlett, Llion Jones, Alexandre Lacoste and izzeddin gur	<i>Accurate Supervised and Semi-Supervised Machine Reading for Long Documents</i>
Cong Duy Vu Hoang, Gholamreza Haffari and Trevor Cohn	<i>Towards Decoding as Continuous Optimisation in Neural Machine Translation</i>
Mark Hopkins, Cristian Petrescu-Prahova, Roie Levin, Ronan Le Bras, Alvaro Herrasti and Vidur Joshi	<i>Beyond Sentential Semantic Parsing: Tackling the Math SAT with a Cascade of Tree Transducers</i>
Tobias Horsmann and Torsten Zesch	<i>Do LSTMs really work so well for PoS tagging? -- A replication study</i>
Nabil Hossain, John Krumm, Lucy Vanderwende, Eric Horvitz and Henry Kautz	<i>Filling the Blanks (hint: plural noun) for Mad Libs Humor</i>
Lifu Huang, Avirup Sil, Heng Ji and Radu Florian	<i>Improving Slot Filling Performance with Attentive Neural Networks on Dependency Structures</i>
Danqing Huang, Shuming Shi, Chin-Yew Lin and Jian Yin	<i>Learning Fine-Grained Expressions to Solve Math Word Problems</i>
Kai Hui, Andrew Yates, Klaus Berberich and Gerard de Melo	<i>PACRR: A Position-Aware Neural IR Model for Relevance Matching</i>
Pierre Isabelle, Colin Cherry and George Foster	<i>A Challenge Set Approach to Evaluating Machine Translation</i>
Masaru Isonuma, Toru Fujino, Junichiro Mori, Yutaka Matsuo and Ichiro Sakata	<i>Extractive Summarization Using Multi-Task Learning with Document Classification</i>
Yangfeng Ji, Chenhao Tan, Sebastian Martschat, Yejin Choi and Noah A. Smith	<i>Dynamic Entity Representations in Neural Language Models</i>
Robin Jia and Percy Liang	<i>Adversarial Examples for Evaluating Reading Comprehension Systems</i>

Authors	Title
Yohan Jo, Michael Yoder, Hyeju Jang and Carolyn Rose	<i>Modeling Dialogue Acts with Content Word Filtering and Speaker Preferences</i>
Kenneth Joseph, Lisa Friedland, William Hobbs, David Lazer and Oren Tsur	<i>ConStance: Modeling Annotation Contexts to Improve Stance Classification</i>
Nobuhiro Kaji and Hayato Kobayashi	<i>Incremental Skip-gram Model with Negative Sampling</i>
Dongyeop Kang, Varun Gangal, Ang Lu, Zheng Chen and Eduard Hovy	<i>Detecting and Explaining Causes From Text For a Time Series Event</i>
Jungo Kasai, Bob Frank, Tom McCoy, Owen Rambow and Alexis Nasr	<i>TAG Parsing with Neural Networks and Vector Representations of Supertags</i>
Katherine Keith, Abram Handler, Michael Pinkham, Cara Magliozzi, Joshua McDuffie and Brendan O'Connor	<i>Identifying civilians killed by police with distantly supervised entity-event extraction</i>
Nikita Kitaev and Dan Klein	<i>Where is Misty? Interpreting Spatial Descriptors by Modeling Regions in Space</i>
Jayant Krishnamurthy, Pradeep Dasigi and Matt Gardner	<i>Neural Semantic Parsing with Type Constraints for Semi-Structured Tables</i>
Guokun Lai, Qizhe Xie, Hanxiao Liu, Yiming Yang and Eduard Hovy	<i>RACE: Large-scale ReAding Comprehension Dataset From Examinations</i>
Wuwei Lan, Siyu Qiu, Hua He and Wei Xu	<i>A Continuously Growing Dataset of Sentential Paraphrases</i>
Man Lan, Jianxiang Wang, Yuanbin Wu, Zheng-Yu Niu and Haifeng Wang	<i>Multi-task Attention-based Neural Networks for Implicit Discourse Relationship Representation and Identification</i>
Carolín Lawrence, Artem Sokolov and Stefan Riezler	<i>Counterfactual Learning from Bandit Feedback under Deterministic Logging : A Case Study in Statistical Machine Translation</i>

Authors	Title
Kenton Lee, Luheng He, Mike Lewis and Luke Zettlemoyer	<i>End-to-end Neural Coreference Resolution</i>
Guang-He Lee and Yun-Nung Chen	<i>MUSE: Modularizing Unsupervised Sense Embeddings</i>
Mike Lewis, Denis Yarats, Yann Dauphin, Devi Parikh and Dhruv Batra	<i>Deal or No Deal? End-to-End Learning of Negotiation Dialogues</i>
Piji Li, Wai Lam, Lidong Bing and Zihao Wang	<i>Deep Recurrent Generative Decoder for Abstractive Text Summarization</i>
Piji Li, Wai Lam, Lidong Bing, Weiwei Guo and Hang Li	<i>Cascaded Attention based Unsupervised Information Distillation for Compressive Summarization</i>
Bofang Li, Tao Liu, Zhe Zhao, Buzhou Tang, Aleksandr Drozd, Anna Rogers and Xiaoyong Du	<i>Investigating Different Syntactic Context Types and Context Representations for Learning Word Embeddings</i>
Jiwei Li and Dan Jurafsky	<i>Neural Net Models of Open-domain Discourse Coherence</i>
Jiwei Li, Will Monroe, Tianlin Shi, S'ebastien Jean, Alan Ritter and Dan Jurafsky	<i>Adversarial Learning for Neural Dialogue Generation</i>
Haoran Li, Junnan Zhu, Cong Ma, Jiajun Zhang and Chengqing Zong	<i>Multi-modal Summarization for Asynchronous Collection of Text, Image, Audio and Video</i>
Hongyu Lin, Le Sun and Xianpei Han	<i>Reasoning with Heterogeneous Knowledge for Commonsense Machine Comprehension</i>
Liyuan Liu, Xiang Ren, Qi Zhu, Shi Zhi, Huan Gui, Heng Ji and Jiawei Han	<i>Heterogeneous Supervision for Relation Extraction: A Representation Learning Approach</i>
Rui Liu, Junjie Hu, Wei Wei, Zi Yang and Eric Nyberg	<i>Structural Embedding of Syntactic Trees for Machine Comprehension</i>
Yang Liu, Kun Han, Zhao Tan and Yun Lei	<i>Using Context Information for Dialog Act Classification in DNN Framework</i>



Authors	Title
Yang Liu and Mirella Lapata	<i>Learning Contextually Informed Representations for Linear-Time Discourse Parsing</i>
Pengfei Liu, Kaiyu Qian, Xipeng Qiu and Xuanjing Huang	<i>Idiom-Aware Compositional Distributed Semantics</i>
Yunfei Long, Lu Qin, Rong Xiang, Minglei Li and Chu-Ren Huang	<i>A Cognition Based Attention Model for Sentiment Analysis</i>
Teng Long, Emmanuel Bengio, Ryan Lowe, Jackie Chi Kit Cheung and Doina Precup	<i>World Knowledge for Reading Comprehension: Rare Entity Prediction with Hierarchical LSTMs Using External Descriptions</i>
Yi Luan, Mari Ostendorf and Hannaneh Hajishirzi	<i>Scientific Information Extraction with Semi-supervised Neural Tagging</i>
Bingfeng Luo, Yansong Feng, Jianbo Xu, Xiang Zhang and Dongyan Zhao	<i>Learning to Predict Charges for Criminal Cases with Legal Basis</i>
Veronica Lynn, Youngseo Son, Vivek Kulkarni, Niranjan Balasubramanian and H. Andrew Schwartz	<i>Human Centered NLP with User-Factor Adaptation</i>
Qingsong Ma, Yvette Graham, Timothy Baldwin and Qun Liu	<i>Further Investigation into Reference Bias in Monolingual Evaluation of Machine Translation</i>
Aroma Mahendru, Viraj Prabhu, Akrit Mohapatra, Dhruv Batra and Stefan Lee	<i>The Promise of Premise: Harnessing Question Premises in Visual Question Answering</i>
Ana Marasovic, Leo Born, Juri Opitz and Anette Frank	<i>A Mention-Ranking Model for Abstract Anaphora Resolution</i>
Diego Marcheggiani and Ivan Titov	<i>Encoding Sentences with Graph Convolutional Networks for Semantic Role Labeling</i>
André F. T. Martins and Julia Kreutzer	<i>Learning What's Easy: Fully Differentiable Neural Easy-First Taggers</i>
Stephen Mayhew, Chen-Tse Tsai and Dan Roth	<i>Cheap Translation for Cross-Lingual Named Entity Recognition</i>

Authors	Title
Lara McConnaughey, Jennifer Dai and David Bamman	<i>The Labeled Segmentation of Printed Books</i>
Dheeraj Mekala, Vivek Gupta, Bhargavi Paranjape and Harish Karnick	<i>SCDV : Sparse Composite Document Vectors using soft clustering over distributional representations</i>
Yuanliang Meng, Anna Rumshisky and Alexey Romanov	<i>Temporal Information Extraction for Question Answering Using Syntactic Dependencies in an LSTM-based Architecture</i>
Dipendra Misra, John Langford and Yoav Artzi	<i>Mapping Instructions and Visual Observations to Actions with Reinforcement Learning</i>
Alex Morales and Chengxiang Zhai	<i>Identifying Humor in Reviews using Background Text Sources</i>
Aldrian Obaja Muis and Wei Lu	<i>Labeling Gaps Between Words: Recognizing Overlapping Mentions with Mention Separators</i>
Ndapandula Nakashole and Raphael Flauger	<i>Knowledge Distillation for Bilingual Dictionary Induction</i>
Shashi Narayan, Claire Gardent, Shay B. Cohen and Anastasia Shimorina	<i>Split and Rephrase</i>
Khanh Nguyen, Hal Daumé III and Jordan Boyd-Graber	<i>Reinforcement Learning for Bandit Neural Machine Translation with Simulated Human Feedback</i>
Kim Anh Nguyen, Maximilian Köper, Sabine Schulte im Walde and Ngoc Thang Vu	<i>Hierarchical Embeddings for Hypernymy Detection and Directionality</i>
Giannis Nikolentzos, Polykarpos Meladianos, Francois Rousseau, Yannis Stavrakas and Michalis Vazirgiannis	<i>Shortest-Path Graph Kernels for Document Similarity</i>
Qiang Ning, Zhili Feng and Dan Roth	<i>A Structured Learning Approach to Temporal Relation Extraction</i>
Rodrigo Nogueira and Kyunghyun Cho	<i>Task-Oriented Query Reformulation with Reinforcement Learning</i>

Authors	Title
Jekaterina Novikova, Ondřej Dušek, Amanda Cercas Curry and Verena Rieser	<i>Why We Need New Evaluation Metrics for NLG</i>
Prakhar Ojha and Partha Talukdar	<i>KGEval: Accuracy Estimation of Automatically Constructed Knowledge Graphs</i>
Deepak P, Dinesh Garg and Shirish Shevade	<i>Latent Space Embedding for Retrieval in Question-Answer Archives</i>
Sungjoon Park, JinYeong Bak and Alice Oh	<i>Rotated Word Vector Representations and their Interpretability</i>
Tommaso Pasini and Roberto Navigli	<i>Train-O-Matic: Large-Scale Supervised Word Sense Disambiguation in Multiple Languages without Manual Training Data</i>
Jasabanta Patro, Bidisha Samanta, Saurabh Singh, Abhipsa Basu, Prithwish Mukherjee, Monojit Choudhury and Animesh Mukherjee	<i>All that is English may be Hindi: Enhancing language identification through automatic ranking of the likeliness of word borrowing in social media</i>
John Pavlopoulos, Prodromos Malakasiotis and Ion Androutsopoulos	<i>Deeper Attention to Abusive User Content Moderation</i>
Haoruo Peng, Ming-Wei Chang and Wen-tau Yih	<i>Maximum Margin Reward Networks for Learning from Explicit and Implicit Supervision</i>
Baolin Peng, Xiujun Li, Lihong Li, Jianfeng Gao, Asli Celikyilmaz, Sungjin Lee and Kam-Fai Wong	<i>Composite Task-Completion Dialogue Policy Learning via Hierarchical Deep Reinforcement Learning</i>
Yuval Pinter, Robert Guthrie and Jacob Eisenstein	<i>Mimicking Word Embeddings using Subword RNNs</i>
Lahari Poddar, Wynne Hsu and Mong Li Lee	<i>Author-aware Aspect Topic Sentiment Model to Retrieve Supporting Opinions from Reviews</i>
Peter Potash and Anna Rumshisky	<i>Towards Debate Automation: a Recurrent Model for Predicting Debate Winners</i>

Authors	Title
Peter Potash, Alexey Romanov and Anna Rumshisky	<i>Here's My Point: Joint Pointer Architecture for Argument Mining</i>
Alessandro Raganato, Claudio Delli Bovi and Roberto Navigli	<i>Neural Sequence Learning Models for Word Sense Disambiguation</i>
Afshin Rahimi, Timothy Baldwin and Trevor Cohn	<i>Continuous Representation of Location for Geolocation and Lexical Dialectology using Mixture Density Networks</i>
Jonathan Raiman and John Miller	<i>Globally Normalized Reader</i>
Prajit Ramachandran, Peter Liu and Quoc Le	<i>Improving sequence to sequence learning with unlabeled data</i>
Farzana Rashid and Eduardo Blanco	<i>Dimensions of Interpersonal Relationships: Corpus and Experiments</i>
Siva Reddy, Oscar Täckström, Slav Petrov, Mark Steedman and Mirella Lapata	<i>Universal Semantic Parsing</i>
Marek Rei, Luana Bulat, Douwe Kiela and Ekaterina Shutova	<i>Grasping the Finer Point: A Supervised Similarity Network for Metaphor Detection</i>
Nils Reimers and Iryna Gurevych	<i>Reporting Score Distributions Makes a Difference: Performance Study of LSTM-networks for Sequence Tagging</i>
Guy D. Rosin, Kira Radinsky and Eytan Adar	<i>Learning Word Relatedness over Time</i>
Sebastian Ruder and Barbara Plank	<i>Learning to select data for transfer learning with Bayesian Optimization</i>
Mrinmaya Sachan, Kumar Dubey and Eric Xing	<i>From Textbooks to Knowledge: A Case Study in Harvesting Axiomatic Knowledge from Textbooks to Solve Geometry Problems</i>
Enrico Santus, Emmanuele Chersoni, Alessandro Lenci and Philippe Blache	<i>Measuring Thematic Fit with Distributional Feature Overlap</i>

Authors	Title
Alexandra Schofield, Laure Thompson and David Mimno	<i>Quantifying the Effects of Text Duplication on Semantic Models</i>
Sarah Schulz and Jonas Kuhn	<i>Multi-modular domain-tailored OCR post-correction</i>
H. Andrew Schwartz, Masoud Rouhizadeh, Michael Bishop, Philip Tetlock, Barbara Mellers and Lyle Ungar	<i>Assessing Objective Recommendation Quality through Political Forecasting</i>
Stanislau Semeniuta, Aliaksei Severyn and Erhardt Barth	<i>A Hybrid Convolutional Variational Autoencoder for Text Generation</i>
Iulian Vlad Serban, Alexander G. Ororbia, Joelle Pineau and Aaron Courville	<i>Piecewise Latent Variables for Neural Variational Text Processing</i>
Yuanlong Shao, Stephan Gouws, Denny Britz, Anna Goldie, Brian Strope and Ray Kurzweil	<i>Generating High-Quality and Informative Conversation Responses with Sequence-to-Sequence Models</i>
Gehui Shen, Yunlun Yang and Zhi-Hong Deng	<i>Inter-Weighted Alignment Network for Sentence Pair Modeling</i>
Tianze Shi, Liang Huang and Lillian Lee	<i>Fast(er) Exact Decoding and Global Training for Transition-Based Dependency Parsing via a Minimal Feature Set</i>
Masumi Shirakawa, Takahiro Hara and Takuya Maekawa	<i>Never Abandon Minorities: Exhaustive Extraction of Bursty Phrases on Microblogs Using Set Cover Problem</i>
Matthias Sperber, Graham Neubig, Jan Niehues and Alex Waibel	<i>Neural Lattice-to-Sequence Models for Uncertain Inputs</i>
Shashank Srivastava, Igor Labutov and Tom Mitchell	<i>Joint Concept Learning and Semantic Parsing from Natural Language Explanations</i>
Felix Stahlberg and Bill Byrne	<i>Unfolding and Shrinking Neural Machine Translation Ensembles</i>
Miloš Stanojević and Raquel Garrido Alhama	<i>Neural Discontinuous Constituency Parsing</i>

Authors	Title
Emma Strubell, Patrick Verga, David Belanger and Andrew McCallum	<i>Fast and Accurate Entity Recognition with Iterated Dilated Convolutions</i>
Tzu-ray Su and Hung-yi Lee	<i>Learning Chinese Word Representations From Glyphs Of Characters</i>
Yu Su and Xifeng Yan	<i>Cross-domain Semantic Parsing via Paraphrasing</i>
Sandesh Swamy, Alan Ritter and Marie-Catherine de Marneffe	<i>"i have a feeling trump will win.....": Forecasting Winners and Losers from User Predictions on Twitter</i>
Chuanqi Tan, Furu Wei, Pengjie Ren, Weifeng Lv and Ming Zhou	<i>Entity Linking for Queries by Searching Wikipedia Sentences</i>
Julien Tissier, Christopher Gravier and Amaury Habrard	<i>Dict2vec : Learning Word Embeddings using Lexical Dictionaries</i>
David Vilares and Yulan He	<i>Detecting Perspectives in Political Debates</i>
Ivan Vulić, Nikola Mrkšić and Anna Korhonen	<i>Cross-Lingual Induction and Transfer of Verb Classes Based on Word Vector Space Specialisation</i>
Henning Wachsmuth, Giovanni Da San Martino, Dora Kiesel and Benno Stein	<i>The Impact of Modeling Overall Argumentation with Tree Kernels</i>
Chuan Wang and Nianwen Xue	<i>Getting the Most out of AMR Parsing</i>
Kexiang Wang, Tianyu Liu, Zhifang Sui and Baobao Chang	<i>Affinity-Preserving Random Walk for Multi-Document Summarization</i>
Di Wang, Nebojsa Jojic, Chris Brockett and Eric Nyberg	<i>Steering Output Style and Topic in Neural Response Generation</i>
Zhongqing Wang and Yue Zhang	<i>Opinion Recommendation Using A Neural Model</i>

Authors	Title
Chengyu Wang, Xiaofeng He and Aoying Zhou	<i>A Short Survey on Taxonomy Learning from Text Corpora: Issues, Resources and Recent Advances</i>
Chengyu Wang, Yan Fan, Xiaofeng He and Aoying Zhou	<i>Learning Fine-grained Relations from Chinese User Generated Categories</i>
Xing Wang, Zhaopeng Tu, Deyi Xiong and Min Zhang	<i>Translating Phrases in Neural Machine Translation</i>
Yasheng Wang, Yang Zhang and Bing Liu	<i>Sentiment Lexicon Expansion Based on Neural PU Learning, Double Dictionary Lookup, and Polarity Association</i>
Leyi Wang and Rui Xia	<i>Sentiment Lexicon Construction with Representation Learning Based on Hierarchical Sentiment Supervision</i>
Zhongqing Wang, Yue Zhang and Ching-Yun Chang	<i>Integrating Order Information and Event Relation for Script Event Prediction</i>
Yan Wang, Xiaojiang Liu and Shuming Shi	<i>Deep Neural Solver for Math Word Problems</i>
Chenguang Wang, Alan Akbik, Iana Chiticariu, Yunyao Li, Fei Xia and Anbang Xu	<i>CROWD-IN-THE-LOOP: A Hybrid Approach for Annotating Semantic Roles</i>
Shaolei Wang, Wanxiang Che, Yue Zhang, Meishan Zhang and Ting Liu	<i>Transition-Based Disfluency Detection using LSTMs</i>
Marlies van der Wees, Arianna Bisazza and Christof Monz	<i>Dynamic Data Selection for Neural Machine Translation</i>
Rongxiang Weng, Shujian Huang, Zaixiang Zheng, XIN-YU DAI and Jiajun CHEN	<i>Neural Machine Translation with Word Predictions</i>
John Wieting, Jonathan Mallinson and Kevin Gimpel	<i>Learning Paraphrastic Sentence Embeddings from Back-Translated Bitext</i>
Derry Tanti Wijaya, Brendan Callahan, John Hewitt, Jie Gao, Xiao Ling, Marianna Apidianaki and Chris Callison-Burch	<i>Learning Translations via Matrix Completion</i>

Authors	Title
Sam Wiseman, Alexander Rush and Stuart Shieber	<i>Challenges in Data-to-Document Generation</i>
Wenhan Xiong, Thien Hoang and William Yang Wang	<i>DeepPath: A Reinforcement Learning Method for Knowledge Graph Reasoning</i>
Zhen Xu, Bingquan Liu, Baoxun Wang, Chengjie SUN, Xiaolong Wang, Zhuoran Wang and Chao Qi	<i>Neural Response Generation via GAN with an Approximate Embedding Layer</i>
Kui Xu and Xiaojun Wan	<i>Towards a Universal Sentiment Classifier in Multiple languages</i>
Hitomi Yanaka, Koji Mineshima, Pascual Martínez-Gómez and Daisuke Bekki	<i>Determining Semantic Textual Similarity using Natural Deduction Proofs</i>
Fan Yang, Arjun Mukherjee and Eduard Dragut	<i>Satirical News Detection and Analysis using Attention Mechanism and Linguistic Features</i>
Diyi Yang, Aaron Halfaker, Robert Kraut and Eduard Hovy	<i>Identifying Semantic Edit Intentions from Revisions in Wikipedia</i>
Zichao Yang, Phil Blunsom, Chris Dyer and Wang Ling	<i>Reference-Aware Language Models</i>
Bishan Yang and Tom Mitchell	<i>A Joint Sequential and Relational Model for Frame-Semantic Parsing</i>
Baosong Yang, Derek F. Wong, Tong Xiao, Lidia S. Chao and Jingbo Zhu	<i>Towards Bidirectional Hierarchical Representations for Attention-based Neural Machine Translation</i>
Helen Yannakoudakis, Marek Rei, Øistein E. Andersen and Zheng Yuan	<i>Neural Sequence-Labelling Models for Grammatical Error Correction</i>
Lili Yao, Yaoyuan Zhang, Yansong Feng, Dongyan Zhao and Rui Yan	<i>Towards Implicit Content-Introducing for Generative Short-Text Conversation Systems</i>
Andrew Yates, Arman Cohan and Nazli Goharian	<i>Depression and Self-Harm Risk Assessment in Online Forums</i>



Authors	Title
Yichun Yin, Yangqiu Song and Ming Zhang	<i>Document-Level Multi-Aspect Sentiment Classification as Machine Comprehension</i>
Xuwan Yin and Vicente Ordonez	<i>Obj2Text: Generating Visually Descriptive Language from Object Layouts</i>
Qingyu Yin, Yu Zhang, Weinan Zhang and Ting Liu	<i>Chinese Zero Pronoun Resolution with Deep Memory Network</i>
Amir Zadeh, Minghai Chen, Soujanya Poria, Erik Cambria and Louis-Philippe Morency	<i>Tensor Fusion Network for Multimodal Sentiment Analysis</i>
Nasser Zalmout and Nizar Habash	<i>Don't Throw Those Morphological Analyzers Away Just Yet: Neural Morphological Disambiguation for Arabic</i>
Rowan Zellers and Yejin Choi	<i>Zero-Shot Activity Recognition with Verb Attribute Induction</i>
Wenyuan Zeng, Yankai Lin, Zhiyuan Liu and Maosong Sun	<i>Incorporating Relation Paths in Neural Relation Extraction</i>
Xiao Zhang, Yong Jiang, Hao Peng, Kewei Tu and Dan Goldwasser	<i>Semi-supervised Structured Prediction with Neural CRF Autoencoder</i>
Xingxing Zhang and Mirella Lapata	<i>Sentence Simplification with Deep Reinforcement Learning</i>
Meishan Zhang, Yue Zhang and Guohong Fu	<i>End-to-End Neural Relation Extraction with Global Optimization</i>
Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli and Christopher D. Manning	<i>Position-aware Attention and Supervised Data Improve Slot Filling</i>
Meng Zhang, Yang Liu, Huanbo Luan and Maosong Sun	<i>Earth Mover's Distance Minimization for Unsupervised Bilingual Lexicon Induction</i>
Yuchen Zhang, Panupong Pasupat and Percy Liang	<i>Macro Grammars and Holistic Triggering for Efficient Semantic Parsing</i>

Authors	Title
Justine Zhang, Arthur Spirling and Cristian Danescu-Niculescu-Mizil	<i>Asking too much? The rhetorical role of questions in political discourse</i>
Jieyu Zhao, Tianlu Wang, Mark Yatskar, Vicente Ordonez and Kai-Wei Chang	<i>Men Also Like Shopping: Reducing Gender Bias Amplification using Corpus-level Constraints</i>
Zhe Zhao, Tao Liu, Shen Li, Bofang Li and Xiaoyong Du	<i>Ngram2vec: Learning Improved Word Representations from Ngram Co-occurrence Statistics</i>
Xiaoqing Zheng	<i>Incremental Graph-based Neural Dependency Parsing</i>
Honglei Zhuang, Chi Wang, Fangbo Tao, Lance Kaplan and Jiawei Han	<i>Identifying Semantically Deviating Outlier Documents</i>

Demo papers

Authors	Title
Abdalghani Abujabal, Rishiraj Saha Roy, Mohamed Yahya and Gerhard Weikum	<i>QUINT: Interpretable Question Answering over Knowledge Bases</i>
Lasha Abzianidze	<i>LangPro: Natural Language Theorem Prover</i>
Alan Akbik and Roland Vollgraf	<i>The Projector: An Interactive Annotation Projection Visualization Tool</i>
Cecilia Ovesdotter Alm, Benjamin Meyers and Emily Prud'hommeaux	<i>An Analysis and Visualization Tool for Case Study Learning of Linguistic Concepts</i>
Franck Dernoncourt, Ji Young Lee and Peter Szolovits	<i>NeuroNER: an easy-to-use program for named-entity recognition based on neural networks</i>

Authors	Title
Tobias Falke and Iryna Gurevych	<i>GraphDocExplore: A Framework for the Experimental Comparison of Graph-based Document Exploration Techniques</i>
Ivan Habernal, Raffael Hannemann, Christian Pollak, Christopher Kamm, Patrick Pauli and Iryna Gurevych	<i>Argotario: Computational Argumentation Meets Serious Games</i>
Chieh-Yang Huang, Tristan Labetoulle, Ting-Hao Huang, Yi-Pei Chen, Hung-Chen Chen, Vallari Srivastava and Lun-Wei Ku	<i>MoodSwipe: A Soft Keyboard that Suggests MessageBased on User-Specified Emotions</i>
Jaesong Lee, Joong-Hwi Shin and Jun-Seok Kim	<i>Interactive Visualization and Manipulation of Attention-based Neural Machine Translation</i>
Ross Mechanic, Dean Fulgoni, Hannah Cutler, Sneha Rajana, Zheyuan Liu, Bradley Jackson, Anne Cocos, Chris Callison-Burch and Marianna Apidianaki	<i>KnowYourNyms? A Game of Semantic Relationships</i>
Alexander Miller, Will Feng, Dhruv Batra, Antoine Bordes, Adam Fisch, Jiasen Lu, Devi Parikh and Jason Weston	<i>ParlAI: A Dialog Research Software Platform</i>
Alexander Panchenko, Fide Marten, Eugen Ruppert, Stefano Faralli, Dmitry Ustalov, Simone Paolo Ponzetto and Chris Biemann	<i>Unsupervised, Knowledge-Free, and Interpretable Word Sense Disambiguation</i>
Simone Papandrea, Alessandro Raganato and Claudio Delli Bovi	<i>SupWSD: A Flexible Toolkit for Supervised Word Sense Disambiguation</i>
Kyle Richardson and Jonas Kuhn	<i>Function Assistant: A Tool for NL Querying of APIs</i>
Ludwig Richter, Johanna Geiß, Andreas Spitz and Michael Gertz	<i>HeidelPlace: An Extensible Framework for Geoparsing</i>
Aaron Sarnat, Vidur Joshi, Cristian Petrescu-Prahova, Alvaro Herrasti, Brandon Stilson and Mark Hopkins	<i>Interactive Visualization for Linguistic Structure</i>
Nathan Schneider and Chuck Wooters	<i>The NLTK FrameNet API: Designing for Discoverability with a Rich Linguistic Resource</i>

Authors	Title
H. Andrew Schwartz, Salvatore Giorgi, Maarten Sap, Patrick Crutchley, Lyle Ungar and Johannes Eichstaedt	<i>DLATK: Differential Language Analysis ToolKit</i>
Ori Shapira, Hadar Ronen, Meni Adler, Yael Amsterdamer, Judit Bar-Ilan and Ido Dagan	<i>Interactive Abstractive Summarization for Event News Tweets</i>
Felix Stahlberg, Eva Hasler, Danielle Saunders and Bill Byrne	<i>SGNMT -- A Flexible NMT Decoding Platform for Quick Prototyping of New Models and Search Strategies</i>
Kohsuke Yanai, Misa Sato, Toshihiko Yanase, Kenzo Kurotsuchi, Yuta Koreeda and Yoshiki Niwa	<i>StruAP: A Tool for Bundling Linguistic Trees through Structure-based Abstract Pattern</i>